AI COCKPIT ENTERPRISE ARCHITECTURE & REFERENCE MANUAL

AI Cockpit — Enterprise vs. Pilot Governance Guide

Complete Technical & Financial Specification for Autonomous Agent Portfolio Governance

Target Audience: Executive Leadership & Client Pitch Teams | **Cockpit Version:** 2.1.0 Enterprise | **Published:** November 2026 | **Status:** Approved Production Reference

1. Dual-Cohort Architecture: Enterprise View vs. Pilot View

The AI Cockpit separates portfolio governance into two synchronized views, accessible via the top-level toggle:

🍼 Pilot View (4 Live Bots)

- **Purpose:** Client proof-of-concept, live conversation demonstrations, and real-time Azure OpenAI telemetry sync.
- **Scope:** Dedicated pilot agents in IT Service Desk, Supply Chain, AP Invoice Processing, and HR.
- **Telemetry:** Live token consumption, per-user execution latency, active employee session logging, and live chat drawer launch.
- **Scale:** $1.2k – $3.5k monthly net value; proof of execution accuracy.

🏢 Enterprise View (1,200 Agents)

- **Purpose:** Demonstrates macro organizational impact, labor capacity unlocking, and FinOps budgeting at Fortune 500 scale.
- **Scope:** 1,200 autonomous agents across 35 teams and 6 business functions (Engineering, Sales, CS, HR, Finance, Operations).
- **Telemetry:** Real-time telemetry for the live pilot cohort + mathematically deterministic telemetry for the remaining fleet.
- **Scale: $18.2M+ Annualized Net Value** (~$1.52M/mo) and **~83,345 hours saved/mo** (~520 FTEs unlocked).

Dimension	Pilot View (4 Live Agents)	Enterprise View (1,200 Agents)
Live Agent Links	Active `[Live Agent ↗]` buttons direct to running container web apps (Copilot Studio, IT Service Desk).	Live links strictly restricted to the 2 demo bots; enterprise fleet displays clean catalog metrics.
Monthly Labor Unlocked	~120 – 350 hours / month (~0.8 – 2.2 FTEs).	**~83,345 hours / month (~520 FTEs capacity unlocked).**
FinOps Cost Tracking	Real Azure OpenAI tokens ($2.50/M input, $10.00/M output) + Copilot Capacity Packs ($200/mo).	Full enterprise breakdown: License seats, token consumption, platform infra, and enablement fees.

2. The 4-Tier Agent Governance Decision Matrix

Every autonomous agent in the cockpit is audited continuously and assigned one of four governance action statuses. Rules evaluate in strict sequence:

Status	Trigger Condition / Parameter	Operational Meaning & Required Action
🧪 **Experiment**	`Days in Production < 60 days`	Agent is newly deployed. Gathering early production baseline telemetry, usage patterns, and user feedback before applying strict ROI thresholds.
🚨 **Intervene**	**Triggered if ANY of these fail:** • Policy Incidents > 0 • Cost vs. Budget > 110% • Eval Pass Rate < 88% • Net Value < $0 (Negative ROI)	Quality degradation, safety violation, budget overrun, or value destruction. Immediate engineering review required: tune golden evaluation datasets, review guardrails, or apply semantic caching.
🚀 **Accelerate**	`Net Value ≥ $2,500/mo` **AND** `Workforce Adoption ≥ 70%` *(with Eval ≥ 88%, Incidents = 0, Cost ≤ 110%)*	Proven high-impact top performer with strong workforce adoption. Candidate for enterprise scaling, multi-region rollout, and increased departmental investment.
⚙️ **Optimize**	Operates safely within normal parameters (Eval ≥ 88%, Cost ≤ 110%, Incidents = 0), but Net Value is under the $2,500 acceleration threshold.	Standard healthy production status. Ongoing continuous tuning: prompt latency optimization, model distillation, and token caching.

3. Financial & Labor Valuation Formulas

1. Hours Saved (Labor Capacity Returned)

```
Hours Saved = Total Successful Runs × Baseline Task Time (hours)
```

Baseline task times represent human benchmark duration: e.g., 20 mins (0.33 hrs) for IT ticket triage; 15 mins (0.25 hrs) for SAP shipment lookup.

2. Gross AI Value ($ USD)

```
Gross Value = Hours Saved × Blended Employee Hourly Rate
```

Blended hourly rate is derived from departmental compensation benchmarks: e.g., $6,500/month ÷ 160 hrs = **$40.63 / hour**.

3. Total AI Run Cost ($ USD)

```
Total Cost = Token Consumption Cost + Monthly License Cost + Platform Infra Cost
```

Token cost reflects real LLM consumption: `(Input Tokens × $2.50/M) + (Output Tokens × $10.00/M)`.

4. Net AI Value ($ USD) & Return on Investment (ROI)

```
Net Value = Gross Labor Value - Total AI Cost
```

```
ROI (%) = (Net Value ÷ Total Cost) × 100%
```

If Net Value is negative, the agent automatically triggers the 🚨 **Intervene** action.

5. FTE Capacity Equivalent (Full-Time Employees Unlocked)

```
FTE Unlocked = Total Hours Saved ÷ 160 hours/month
```

Represents employee capacity unlocked from repetitive manual operations to focus on core strategic initiatives.

Key Takeaway for Client Pitches: The AI Cockpit provides end-to-end visibility from a single LLM token prompt up to enterprise-wide FinOps governance ($18M+ net value), ensuring clients can safely deploy, monitor, and scale AI agents with rigorous financial accountability.