WCT // AI PERFORMANCE & GOVERNANCE COCKPIT

Enterprise Telemetry Control Plane · Financial Attribution & User Guide

Live Executive Dashboard: https://wct-cockpit.kindbay-88317352.eastus2.azurecontainerapps.io

Enterprise Scope: WaferWire Cloud Technologies

---

1. The Enterprise Executive Challenge

Enterprise leaders face a critical challenge: ***'We are investing heavily across AI licenses, custom LLM agents, and developer tools, but what verifiable business value and productivity are we actually getting back?'***

When leadership asks for 'AI Productivity', they require **4 verifiable financial and operational proofs** rather than raw vanity metrics:

1. Labor Hours Saved Exact human work-hours returned to the workforce by automating manual workflows.	**2. Gross Labor Value** Dollar worth of saved time based on departmental wage baselines ($/hr).	**3. Total AI Cost Governance** Complete visibility into fixed seats, Copilot packs, and cloud LLM tokens.	**4. Net AI Value & ROI** Bottom-line net financial return after subtracting 100% of licenses and compute.

2. The 4 Enterprise AI Operational Tiers

To govern enterprise AI accurately, you cannot treat all AI the same. Across the Microsoft enterprise ecosystem, the Cockpit orchestrates AI across 4 distinct operational tiers:

Operational Tier	Architecture & Pricing Model	Configured Inputs (Admin & FinOps)	Automated Live Telemetry (Cockpit)
Tier 1: Microsoft 365 Copilot (Teams, Chat, Word, Excel)	Fixed Seat Subscription $30.00 / user / month flat	• Total active purchased licenses • Knowledge Worker wage baseline • 5 min/prompt (MSFT Research)	• Real WaferWire telemetry stream • Cumulative human hours returned • Automatic dormant seat detection
Tier 2: Autonomous Copilot Studio (Global Supply Chain Bot)	Low-Code Multi-Agent Workflow $200/mo Pack (25k queries)	• Baseline manual task handling time • Logistics Specialist wage baseline • Option A (App Insights) / Option B	• Interactive live Teams chat interface • 4 Gateways (SAP, FedEx, SNOW) • Billed queries & SLA pass rate tracking
Tier 3: Pro-Code Foundation Agents (IT Ticket Triage Agent)	Custom Azure OpenAI Python Bot Metered Token Economics	• Manual triage baseline (minutes) • Helpdesk Technician wage rate • Token rates: $2.50 in / $10.00 out per 1M	• Fine-grained prompt & completion tokens • SLA Evaluation Pass Rate verification • Live interactive ticket routing demo
Tier 4: Developer AI Tooling (Antigravity / IDE Tools)	Developer Seat Productivity $19 - $30 / developer / month	• Software Engineer salary baseline • Tool license cost per developer	• Real developer team coding telemetry • Suggestion acceptance rate tracking • Engineering hours saved on boilerplate

3. Executive Governance & Decision Matrix

The Cockpit portfolio engine continuously evaluates each agent against authorized thresholds and categorizes it into one of 4 governance buckets: **ACCELERATE** (high ROI & high eval accuracy — scale deployment company-wide); **OPTIMIZE** (positive value but requires prompt/token tuning); **INTERVENE** (negative ROI or SLA eval failure — halt rollout); **EXPERIMENT** (early pilot <30 days in production).

INTERACTIVE PLATFORM TOUR (TABS 1 TO 4)

Live Feature Exploration Guide · How to Test Live Agents · In-Drawer Telemetry Sync

1. Executive Summary & Agent Scorecard (Tabs 1 & 2)

• **Tab 1 — Executive Summary:** C-Suite macro command center displaying total company **Net AI Value Created**, **Total Hours Saved**, **Workforce Adoption %**, and **Total AI Cost**. Use the top Department dropdown to filter rollups by Engineering & IT, Finance, or Operations.
• **Tab 2 — Agent Scorecard (Interactive Heatmap & 2 Live Testing Agents):** Organizes working agents across business functions with live interactive testing available directly in the browser:

Live Agent	Platform & Gateways	How to Test Live in Browser	In-Drawer Real-Time Refresh
Global Supply Chain Bot (Operations · Supply Chain)	• Microsoft Copilot Studio • SAP S/4HANA (BAPI) • FedEx Freight API • ServiceNow REST & Dataverse	1. Click the blue **[Live Agent]** button. 2. Interactive Teams interface opens. 3. Test *'Track PO-98412'* or *'Issue RMA'*. 4. Real adaptive card response delivered.	In the inspection drawer, click **[Refresh Metrics]** next to *Open Live Agent*. The Cockpit instantly updates query count, handling time saved, and cost in-place **without reloading the website**!
Ticket Triage Agent (Eng & IT · IT Service Desk)	• Azure OpenAI GPT-4o • IT Helpdesk Routing • Metered token economics	1. Click the **Ticket Triage Agent** row. 2. Right-hand drawer slides out. 3. Click **[Open Live Agent]**. 4. Submit an IT issue; agent classifies priority, team, and SLA in <1.2s.	Click **[Refresh Metrics]** in the drawer header. Daily employee usage, prompt/completion tokens, and user sessions update live from the database!

2. FinOps Governance & Employee Adoption (Tabs 3 & 4)

• **Tab 3 — Cost & Budget:** Dedicated FinOps command center showing total spend safely within authorized budget ceilings. Features Departmental Cost Allocation Treemaps, 30-day run-rate forecasting, and cost composition breakdown (Fixed Seats, Capacity Packs, and Variable LLM Tokens) to eliminate unexpected cloud billing spikes.
• **Tab 4 — Seats & Adoption:** Aggregates authentic WaferWire employee telemetry. To comply with corporate privacy guidelines, employee names and corporate emails are anonymized via **synthetic identity masking** (e.g., *Employee 1, emp1@company.com*), preserving 100% of the true prompt interaction volumes, app footprints (Teams, Word, Excel), and timestamps without exposing PII. Features automated **Dormant Seat Identification** (detecting idle licenses to save subscription spend) and **Developer AI Tooling (Antigravity)** tracking engineering sessions and code suggestion acceptance rates.

ADMIN & CONFIG PORTAL (TAB 5) & VALUATION SANDBOX

Super Admin Operational Sub-Sections · Onboarding Architecture · Dynamic Mathematical Proofs

1. Admin & Config: Detailed Breakdown of the 8 Sub-Sections

The **Admin & Config** portal is restricted by Role-Based Access Control (RBAC) to Super Administrators. It gives enterprise administrators complete operational control across 8 dedicated governance sub-sections:

Admin Sub-Section	Operational Purpose & Capabilities
1. AI Subscriptions & Licenses	Defines baseline seat license pricing for Microsoft 365 Copilot ($30/mo), GitHub Copilot ($19/mo), and Copilot Studio Capacity Packs ($200/mo flat for 25k queries) to anchor recurring fixed costs.
2. Labor Rates & Productivity	Configures hourly employee wage benchmarks by business function (Engineering, Operations, Finance, Sales) and standard monthly FTE working hours (160h) used to quantify hours saved into Gross Value.
3. FinOps Quotas & Guardrails	Establishes enterprise budget variance thresholds, maximum monthly cloud token spend caps, and automated drift alerts to enforce fiscal discipline and prevent surprise billing overruns.
4. Governance Action Rules	Sets policy thresholds (Net AI Value targets, adoption rates, SLA pass rate limits) that automatically categorize agents into ACCELERATE, OPTIMIZE, INTERVENE, or EXPERIMENT.
5. Telemetry Ingestion (CSV / Excel)	Provides a bulk-import utility to upload offline telemetry spreadsheets and historical usage exports directly into the Cockpit database when direct API streaming is not yet connected.
6. Developer AI Tools	Configures engineering IDE tool parameters (Antigravity, GitHub Copilot) including seats assigned, developer hourly wage rates, and velocity benchmarks to measure software development ROI.
7. GEMINI / COPILOT	Enables side-by-side comparative benchmarking parameters between Google Gemini Enterprise and Microsoft Copilot ecosystems for cost-per-prompt, response latency, and user adoption.
8. Agents Onboarding	Central registry to onboard enterprise agents, define infrastructure fees, token economics (input/output per 1M), baseline task minutes, and connect telemetry via Option A (Azure App Insights) or Option B (Dataverse OData).

2. Mathematical Proof of the Live Valuation Sandbox

In the Admin & Config portal, leadership can review the **Live Portfolio Valuation Impact Sandbox**. Every metric dynamically recalculates in real-time as telemetry events stream into the database:

Metric Name	Executive Meaning	Exact Enterprise Formula	Live Valuation Attribution Logic
Blended Labor Rate	Enterprise Average Wage	Weighted across active functions: (Eng 40%, Ops 30%, Fin 15%, Sales 15%)	Combines departmental salaries into a unified enterprise baseline ($/hr) for valuing automated hours.
Gross Labor Value	Total Productive Value	Total Hours Saved × Blended Wage Rate	Direct monetary worth of all human work-hours returned to the company by active agents.
Net AI Value	Bottom-Line Net Profit	Gross Labor Value – Total AI Cost	True net value generated after subtracting 100% of licenses, capacity packs, and token fees.
ROI Multiplier	Investment Return Ratio	Gross Labor Value ÷ Total AI Cost	Quantifies total business return generated per $1.00 spent on AI compute, software, and seats.
Labor Capacity Unlocked	FTE Work Force Equivalent	Total Hours Saved ÷ Monthly FTE Hours (160h)	Translates hours saved into equivalent full-time knowledge workers added to corporate capacity.